

# **When Knowing Is Not Enough: Designing Pre-Execution Governance for LLM Agents in Organizational Systems**

## **Abstract**

Organizations increasingly deploy large language models (LLMs) as execution agents in high-stakes information systems (IS) workflows such as financial analysis, compliance monitoring, and automated reporting. A recurrent failure mode has emerged: LLM agents accurately diagnose that their available context is defective, yet proceed to act on it when delegation is unconstrained. This compromises auditability and exposes organizations to compliance and reputational risk. Current AI governance frameworks specify policy and model-level controls but leave the execution layer of LLM-agent pipelines largely unspecified. Drawing on principal–agent theory as the main governance lens, with information design and construal level theory as supporting perspectives, we design GateFix: a pre-execution governance mechanism operationalized through a four-dimensional context quality (4D-CQ) check and a deterministic authorization rule. We evaluate it in a within-subject experiment on 500 financial reasoning tasks across four commercial LLMs and three benchmark corpora. For models above a laboratory-derived 70% diagnostic-accuracy floor, GateFix reduces false authorization on defective contexts from roughly 50% to under 5%. Safety gains transfer across task domains, while accuracy gains attenuate as task abstraction rises. The paper extends IS/AI governance theory to the execution layer by identifying structural routing failure as a new governance failure mode, delivers a deployable execution-harness artifact that organizations can embed in agentic pipelines, including ReAct and RAG agent loops, and derives three design principles for AI-enabled execution systems: safety–accuracy separability, minimum diagnostic capability, and task-structure contingency.

Keywords: AI governance; LLM agents; agentic AI systems; execution-layer governance; context quality; principal–agent theory; design science research

## 1. Introduction

Consider a senior analyst computing year-over-year revenue growth. She flags two problems: the primary 10-K source is missing the 2018 revenue figure, and secondary sources disagree by over four billion dollars. She documents both discrepancies—and then executes the calculation anyway. The result carries an 8.1 percentage-point error relative to the audited ground truth. When organizations push this kind of decision onto an LLM agent operating inside an automated pipeline, the same failure mode re-emerges at scale, and the organization inherits the risk. The scale-up matters: a human analyst makes this mistake once per report. An LLM agent makes it at pipeline throughput, propagating defective outputs into downstream financial, operational, and compliance systems before any human reviewer sees them.

Figure 1 illustrates this governance gap and the proposed response. Panel A depicts what we later formally define as a structural routing failure: even when an LLM agent accurately assesses that context quality is defective, its execution pathway is not structurally bound to that assessment—diagnosis and authorization are decoupled. Panel B shows the consequence of ungoverned execution: errors propagate downstream into organizational information flows, with no gatekeeping at the execution boundary. Panel C shows the governed alternative: GateFix interposes a structured diagnostic stage between the agent's quality assessment and its execution decision, converting it into a constrained relationship where defect detection triggers refusal or clarification before any computation occurs.




Figure 1. The governance gap and GateFix response. Panel A — Structural routing failure: even when the LLM agent accurately assesses that context quality is defective, the diagnostic signal is not routed into its execution authorization, so the agent proceeds anyway. Panel B — Consequence of ungoverned execution: a single defective output propagates downstream into multiple organizational systems (financial reporting and risk, operational decision systems, compliance and audit) before any human reviewer sees it. Panel C — GateFix governed alternative: four binary 4D-CQ checks (Relevance, Coverage, Ordering, Robustness) feed a deterministic authorization rule that issues PASS, CLARIFY, or REFUSE before any computation occurs. Managerial reading: today's agentic pipelines let LLMs act on defective context; GateFix inserts a deterministic checkpoint between what the agent knows and what the agent is authorized to do.

Real-world deployment incidents confirm this structural absence extends beyond individual decision-makers. In August 2024, indirect prompt injection attacks exploited Slack AI's retrieval pipeline to exfiltrate private channel data—the executing model acted on adversarially embedded context with no

pre-execution integrity check [1]; the June 2025 Microsoft 365 Copilot EchoLeak disclosure (CVE-2025-32711)—a zero-click AI command injection vulnerability—followed the same structural pattern [2]. Both incidents share the same structural signature: consequential execution proceeded without evaluation of context fitness at the pre-execution boundary.

Three strands of literature have approached pieces of this problem without converging on the execution layer. First, AI/IT governance research has developed rich accounts of structural, process, and relational governance mechanisms inside organizations [3, 4, 5]; this work has rarely examined the LLM-agent pipeline itself, where the agent's own diagnostic signal about context quality is not structurally linked to its execution authorization. Second, human–AI collaboration research has documented systematic over-reliance on AI recommendations and has studied cognitive forcing functions as a mitigation [6]; this literature focuses on human trust and adoption rather than on how the agent's internal execution decision is made. Third, the emerging literature on LLM delegation and agentic systems has shown that LLM agents require purpose-built governance mechanisms at the contract layer [7, 8, 9]; it has not yet characterized the control points inside the execution pipeline itself.

We address a gap at the intersection of these strands: LLM agents in execution roles often proceed even after accurately diagnosing that their context is defective. We ask two research questions. (RQ1) How can organizations design LLM execution pipelines so that context quality diagnostics structurally govern execution authorization? (RQ2) Under what conditions—across tasks, models, and deployment domains—does such an execution-layer governance mechanism remain effective? We term this class of failures structural routing failure: the agent accurately diagnoses defective context, yet its execution pathway is not structurally constrained by that diagnostic signal. We answer both questions through the design and empirical test of GateFix, a pre-execution governance mechanism operationalized as a four-dimensional context quality check plus a deterministic authorization rule, evaluated in a within-subject experiment across four commercial LLMs on the FinQA financial reasoning benchmark [10] and two transfer corpora [11, 12].

**This paper makes three contributions:**

1. Theoretical contribution. We extend the IS / AI governance conversation from the policy and contract layers to the execution layer, identifying a new governance failure mode—structural routing failure—in which an agent's accurate diagnostic signal is not structurally bound to its execution authorization. This repositions execution-layer pipeline design as a first-class IS governance question alongside contract-layer incentive design and model-layer alignment.

2. Design and practice contribution. We design and empirically validate a deployable execution-layer governance mechanism (GateFix) that organizations can insert into existing LLM-agent pipelines as an authorization micro-service, without retraining. In our experiments it reduces false authorization on defective contexts from approximately 50% to under 5% for capable models, converting previously unreliable agents into operationally trustworthy components. Beyond the mechanism itself, we offer a reusable experimental paradigm—within-subject, multi-model, with a synthetic defect corpus—and three design principles for AI-enabled execution systems. First, safety–accuracy separability: safety and accuracy have different portability characteristics and can be governed separately. Second, minimum diagnostic capability: the mechanism is reliable only above a laboratory-derived candidate threshold of approximately 70% diagnostic accuracy. Third, task-structure contingency: effectiveness is moderated by construal level, so the mechanism delivers the most value on structured, verifiable tasks.

Scope and boundary. This study measures perception–decision decoupling at the level of LLM judgment under controlled conditions, not in live agentic deployments with tool calls or environmental feedback.

The 70% diagnostic capability threshold is a laboratory-derived candidate criterion requiring live-system validation before adoption as an operational standard. We position execution-layer harness design as an IS governance problem: the locus of control, accountability, and auditability for LLM-agent execution lies with organizational information systems, not with model training pipelines.

The remainder of this paper is organized as follows. Section 2 develops the theoretical framework and hypotheses. Section 3 describes the research method. Section 4 presents results. Section 5 discusses theoretical and practical implications. Section 6 addresses limitations, and Section 7 concludes.

## **2. Theoretical Background and Hypotheses**

We develop the theoretical framework in four steps. Section 2.1 locates the execution-layer gap within existing AI/IT governance scholarship and current AI governance frameworks. Section 2.2 uses principal–agent theory as the governance lens and introduces structural routing failure as the specific failure mode we address. Section 2.3 brings in construal level theory as a boundary condition on intervention effectiveness, and Section 2.4 develops context quality and information design as the basis for the 4D-CQ operationalization. Section 2.5 then states the three hypotheses.

### **2.1. AI/IT Governance and the Execution-Layer Gap**

IT governance research has long distinguished between structural, process, and relational mechanisms that organizations use to direct the development and use of information systems [3, 4, 5]; this literature has rarely opened the black box of LLM-agent execution pipelines, where the agent's own diagnostic signal about context quality does not structurally govern its authorization decision. Structural mechanisms—committees, decision rights, and authorization points—separate policy specification from operational enforcement so that execution is bound by rules rather than by individual discretion. As organizations extend IT into AI-enabled execution, these same governance primitives become directly relevant. Yet most AI governance scholarship to date has stayed at the level of high-level principles, policies, and compliance frameworks; the internal pipeline of the LLM agent—where the authorization decision is in fact made—has received much less attention. A growing body of empirical evidence shows why the execution-pipeline layer matters.

Chen et al. [13] provide foundational evidence showing that LLMs replicate well-documented human cognitive biases—anchoring, framing effects, and status quo bias—indicating that LLM outputs are shaped by the information environment, not merely by the model's internal knowledge. Governing LLM execution therefore requires governing the prompt structure, not only the model itself.

Subsequent work extends this insight along multiple dimensions. Kirshner et al. [14] report that task difficulty systematically shapes LLM behavioral patterns, with simple tasks eliciting prosocial behavior and complex tasks eliciting self-interested behavior—suggesting that the abstraction level at which tasks are represented modulates LLM execution tendencies. Calibration research further shows that perceptual accuracy and decisional reliability diverge in neural networks [15], indicating that accurate perception does not guarantee reliable decision-making even in well-trained models. Adversarial evaluations confirm that unrestricted LLM execution produces unauthorized compliance and output misreporting in live environments [16]. Position bias studies document that LLMs exhibit systematic sensitivity to information ordering [17]—a representational vulnerability with direct governance consequences wherever prerequisite quantities must precede dependent computations.

These findings collectively indicate that LLM behavioral biases in execution contexts are (a) systematic across models and model families, (b) shaped by the signal structure rather than fixed model properties, and (c) not self-correcting through capability improvements alone. A governance intervention that restructures the information environment is therefore the theoretically appropriate response to this class of execution failures.

These behavioral properties have a structural consequence for LLM-agent architecture. Current agentic designs—ReAct [18], AutoGen, and related systems—couple perception and execution sequentially but do not make the agent's perception a gating input to its action, producing an open-loop control configuration [19]. Industry risk reports confirm that this is a category-level deployment property [20, 21] rather than an artifact of a particular vendor or model, and it is the specific gap we address at the execution layer.

### **Relation to existing AI governance frameworks.**

The execution-layer gap is not an oversight of existing AI governance frameworks so much as a consequence of their layer of analysis. ISO/IEC 42001, the EU AI Act, and the NIST AI Risk Management Framework all specify governance at the level of organizational policy, risk classification, documentation, and monitoring obligations; industry best-practice guidance (e.g., OWASP Top 10 for

LLM Applications [21]) adds catalogues of failure modes. None of these frameworks specifies how an LLM-agent pipeline should be architected so that the agent's own diagnostic signal about context quality binds its execution decision. Our work is therefore complementary rather than competing. It provides an execution-layer instantiation of these frameworks, specifying a concrete, auditable design pattern by which high-level principles—context integrity, safe execution, and human oversight—are operationalized inside the pipeline. The PASS/CLARIFY/REFUSE log serves as the enforcement artifact that policy-level frameworks require organizations to document and monitor.

### **Execution harness as an IS governance artifact.**

The engineering term execution harness [22] has emerged in frontier AI deployments to describe the governance environment that constrains and directs LLM execution. We adopt this term but re-anchor it theoretically: an execution harness is an IS governance artifact in the Hevner [23] sense, whose design translates organizational rules, constituted at the policy layer, into deterministic pre-execution constraints at the inference layer. Formalizing the harness this way supports progress toward auditable harness design, cross-institutional deployment standards, and a theoretical basis for regulatory oversight—objectives that complement, rather than substitute for, the engineering practice from which the harness concept emerged. GateFix is the concrete harness artifact we study in this paper; the design principles we derive in Section 5 are intended to generalize beyond this particular instantiation.

## **2.2. Principal-Agent Theory as a Governance Lens**

Principal-agent theory [24, 25] governs delegation relationships in which an agent acts on a principal's behalf under conditions of information asymmetry. The framework identifies two canonical failure modes: moral hazard (the agent takes unobservable actions that diverge from the principal's interest) and adverse selection (the agent holds private information the principal cannot verify prior to contracting). A natural question is whether PA theory remains the appropriate analytical lens when the agent is a stateless language model lacking strategic self-interest. Our use of PA theory is not premised on LLMs being utility-maximizing agents in the classical sense; rather, the framework provides the most precise



vocabulary available for describing the structural failure we study—a delegated capacity whose output is not bound to the authorization of action. This usage is consistent with established applications of agency theory in IS, where the framework has been productively extended to non-strategic settings, including IT outsourcing, IS audit, and electronic-market design, without invoking deliberate concealment as a precondition (e.g., Eisenhardt [24]). Alternative lenses are less apt for distinct reasons. Control theory characterizes feedback loops and corrective response but does not address the structural binding of signal to authorization. Mechanism design addresses incentive alignment but presupposes strategic agents whose preferences must be elicited; our problem concerns the routing of an already-produced diagnostic signal. PA theory's information-asymmetry frame, suitably reinterpreted from strategic concealment to structural inaccessibility, remains the closest available theoretical language for this class of governance failures. Applying PA theory to LLM execution requires structural extension because neither canonical failure mode fully describes the governance problem we study. Moral hazard presupposes stable self-directed interests that diverge from the principal's under imperfect monitoring; LLM execution failures arise from structural properties of the processing pipeline, not strategic non-compliance. Adverse selection presupposes pre-contractual asymmetry where the agent holds private information; for LLMs, the relevant diagnostic information is accessible upon explicit query, as our diagnostic prompts demonstrate. What the data reveal (Table III) is a distinctive principal–agent failure. We do not claim the open-loop–versus–closed-loop pattern is novel as a general phenomenon; the decoupling of perception from action is a long-standing concern in cybernetics [19]. The contribution lies in its instantiation at the LLM execution layer and in the information-design remedy framed under principal–agent theory. We formally define a structural routing failure as a pipeline-level governance failure in which (i) an LLM agent possesses accurate diagnostic capacity regarding contextual defects, (ii) this diagnostic signal is accessible to the principal upon query, yet (iii) the execution pathway does not, by default, route the agent's authorization decision through that diagnostic assessment. This is precisely the sense in which knowing is not enough: the agent's accurate diagnostic signal does not, by itself, constrain its behavior unless the pipeline architecture forces that signal into the authorization path. The failure is architectural rather than

informational. Classical moral hazard and adverse selection are insufficient labels: the agent's interests are not divergent and the information is not hidden. The relevant remedy therefore operates through information design—restructuring the signal environment so that the agent's diagnostic output is structurally bound to its execution decision—rather than through incentive realignment or enhanced monitoring.

Boundary of the construct. Structural routing failure is intentionally narrow. It presupposes that the agent already has non-trivial diagnostic capability. The construct is therefore analytically distinct from three other failure modes: data-quality failures, where the upstream signal itself is inadequate; model-alignment or reward-hacking failures, where the model is unwilling to flag defects even when asked; and classical human-incentive PA problems, where the principal's contract design misaligns agent incentives. The failure we study is the structural coupling—or, in current agentic designs, the lack of coupling—between the agent's diagnostic output and its authorization to act, holding diagnostic capacity and incentive structure fixed.

At the contract layer, recent work finds that LLM delegation requires purpose-built governance mechanisms [7] and that agent information structure significantly affects bargaining outcomes [8]. Wu et al. [9] further show that AI agents in delegation-like settings can develop cooperative behavior through interactive learning, illustrating that LLM-based delegation systems exhibit emergent behavioral dynamics that organizational governance must accommodate rather than assume away.

Our study operates at a different layer of the same PA framework: where prior work addresses contract-level incentive design [7, 8], we address execution-level information constraints. These layers are complementary and jointly necessary. Contract-layer governance specifies the delegation relationship and how compliance is verified (the *ex ante* structure). Execution-layer governance ensures that the agent's behavior at the moment of action is structurally constrained by its own diagnostic output (the *ex tempore* structure). Together, they constitute a multi-layer governance architecture for LLM-based execution systems. In this sense, we extend principal–agent theory in IS from incentive misalignment to

architectural misrouting of diagnostic information, specifying a third category of governance failure relevant to AI execution systems.

### **2.3. Task Structure as a Boundary Condition**

Construal level theory (CLT) [26] posits that the level of abstraction at which an event is represented determines the processing mode applied to it. Concrete representations elicit feature-specific, detail-sensitive processing; abstract representations elicit schematic, gist-based processing. Although CLT was developed to explain human cognition, its core prediction—that task structure concreteness modulates judgment quality—generates testable implications for LLM behavior. Empirical evidence that LLMs respond to task difficulty [14] and information environment structure [13] in CLT-consistent ways supports the application of CLT as a predictive framework for LLM governance design, without requiring that LLMs possess the same cognitive mechanisms as humans.

We apply CLT to explain why intervention effects vary with task structure concreteness. Pre-execution information constraints operate through dimensional decomposition: they break an ambiguous execution task into four concrete binary diagnostic subtasks (Is context relevant? Is numerical coverage sufficient? Is ordering coherent? Is the context robust to noise?). This decomposition increases the concreteness of the assessment task, directing the LLM toward feature-by-feature verification rather than holistic processing that may overlook specific defects.

Consistent with CLT, we expect a moderation effect: governance through decomposition will be most effective for tasks that are structurally concrete—tasks with verifiable ground truth, structured data, and well-defined correctness criteria. As tasks become more abstract (open-ended, multi-turn, requiring contextual inference), decomposition loses traction because the diagnostic dimensions themselves become harder to evaluate in binary terms.

This prediction has direct empirical implications. FinQA [10] instances are highly structured single-turn financial reasoning tasks with verifiable numerical answers—low abstraction. TAT-QA [11] introduces hybrid tabular-textual formats requiring cross-format inference—moderate abstraction. ConvFinQA [12]

uses multi-turn conversational reasoning—higher abstraction. Consistent with CLT, we expect intervention effects to attenuate as tasks move from low to high construal level, a prediction we test directly in Section 4.

The finding that task difficulty shapes LLM behavior [14] provides converging evidence, as task difficulty and task concreteness covary.

## 2.4. Context Quality and Information Design

Information design theory [27] studies how a designer can structure the information environment to influence the decisions of agents who observe signals and take actions. In Bergemann and Morris's [27] unified framework, the designer controls the signal structure—what information is revealed, in what form, and with what precision—while agents optimize given the signals they receive. The framework has been applied in mechanism design and persuasion, but has not been extended to the governance of AI execution agents.

We operationalize pre-execution governance as an information design problem. The principal cannot directly control the LLM agent's internal processing, but can control the signal structure presented before execution. Three information design strategies are implemented through the 4D-CQ framework:

1. **Task Decomposition:** The multi-dimensional diagnostic prompt decomposes a holistic quality judgment into four concrete binary assessments (Relevance, Coverage, Ordering, Robustness). This reduces the agent's decision problem from a complex multi-attribute evaluation to a set of verifiable binary signals, lowering construal level and reducing judgment noise. In information design terms, this is a signal refinement that increases the informativeness of the agent's assessment.
2. **Constraint Anchoring:** Each diagnostic dimension specifies explicit decision criteria (e.g., "DEFECT if any key quantity required for the computation is absent, contradictory across sources by more than 5%, or relies on estimates rather than reported figures"). These anchors constrain the agent's inference space, converting an open-ended quality assessment into a bounded classification problem analogous to monotonicity constraints in mechanism design [27].

3. **Role Specification:** The system prompt assigns the agent a specific evaluator role ("You are a financial context quality evaluator"), displacing the default execution-oriented processing mode. This reduces asymmetry by making governance intent explicit.

Table I maps each strategy to its information design foundation, its 4D-CQ operationalization, and the CLT mechanism through which it influences LLM behavior.

**TABLE I: Information Design Strategies in GateFix**

| Strategy             | Information Design<br>Foundation                                                     | 4D-CQ<br>Operationalization                                                       | CLT Mechanism                                                                                           |
|----------------------|--------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|
| Task Decomposition   | Signal refinement: increasing informativeness of agent assessment [27]               | Four binary diagnostic prompts (Relevance, Coverage, Ordering, Robustness)        | Lowers construal level from holistic quality judgment to feature-by-feature verification                |
| Constraint Anchoring | Boundary conditions analogous to monotonicity constraints [27]                       | Explicit numerical thresholds per dimension (e.g., $\leq 5\%$ source discrepancy) | Constrains inference space; prevents high-level gist processing from overriding specific defect signals |
| Role Specification   | Type disclosure: reducing information asymmetry by making governance intent explicit | System prompt assigns "financial context quality evaluator" role                  | Displaces default execution-oriented construal with evaluation-oriented construal                       |

The deterministic authorization function then converts the agent's constrained diagnostic outputs into execution decisions: PASS if all dimensions are adequate; REFUSE if any critical dimension (Relevance, Coverage) fails; CLARIFY if only non-critical dimensions (Ordering, Robustness) fail. This three-state design maps onto organizational escalation workflows, with CLARIFY outcomes routable to human-in-

the-loop review without full task abandonment—preserving throughput while maintaining the safety guarantee.

IS governance scholarship [3, 4] has established the principle of separating policy from mechanism in organizational IT systems; GateFix instantiates this separation at the inference layer, making the authorization policy auditable, configurable, and independent of model modification.

### **Positioning within design science research.**

The study is positioned within the design science research (DSR) tradition in information systems [23, 28]. In DSR terms, GateFix is an instantiated artifact, the 4D-CQ framework and the deterministic authorization rule are its design components, and the three design principles derived in Section 5 (safety–accuracy separability, minimum diagnostic capability, task-structure contingency) constitute the design knowledge this paper contributes. Gregor and Hevner's [28] positioning categories place the contribution at both the artifact level (level 1) and nascent design theory (level 2). Consistent with Hevner et al.'s [23] guidelines, managerial relevance is established by the high-stakes workflows we target, rigor by the theory grounding and within-subject multi-model protocol, and design evaluation by the quantitative cross-model, cross-corpus, and ablation analyses; the primary mode of artifact evaluation is controlled experiment rather than field intervention.

Theoretical map. The four streams above are organized hierarchically. Principal–agent theory and information design serve as the primary backbone: PA theory supplies the language of structural routing failure, and information design specifies why dimensional decomposition plus a deterministic rule is the appropriate remedy. Construal level theory enters as a moderator that explains the task-structure boundary on intervention effectiveness. IT governance and design science research provide the IS context in which the harness is positioned as a governance artifact, but do not by themselves drive the hypotheses tested below.

## **2.5. Hypotheses Development**

The theoretical framework generates three empirically testable hypotheses:

**Hypothesis 1 (Main Effect).** Pre-execution information constraints significantly reduce behavioral biases in LLM execution agents, increasing authorization accuracy relative to unconstrained delegation.

Rationale: PA theory, information design, and CLT jointly predict that constraining the delegation relationship through structured diagnostic decomposition aligns the agent's behavior with its capability: reduced asymmetry improves compliance, constrained signals direct inference, and low-level construal enhances defect detection. Managerially, H1 predicts that inserting a structured pre-execution gate should materially improve whether LLM agents 'do the right thing' conditional on what they already know.

**Hypothesis 2 (CLT Moderation).** Governance effectiveness increases as task construal level decreases—that is, pre-execution information constraints are more effective for concrete, structured tasks than for abstract, open-ended tasks.

Rationale: Consistent with CLT, dimensional decomposition is most effective when the decomposed subtasks are amenable to concrete construal. For well-structured tasks with verifiable numerical answers, the 4D diagnostic prompts map directly onto verifiable features. For tasks requiring cross-format or multi-turn reasoning, the diagnostic dimensions themselves become harder to evaluate concretely, reducing the constraining power of the decomposition strategy. Managerially, H2 predicts that execution-layer governance will deliver the clearest value on structured, verifiable workflows (e.g., financial computation, compliance verification) and smaller gains on abstract, judgment-heavy tasks.

**Hypothesis 3 (Decomposition Granularity).** Finer-grained dimensional decomposition yields monotonically greater bias correction than coarser decomposition, and this ordering is preserved as a structural property of the signal design independent of the specific task corpus.

Rationale: Information design theory predicts that signal precision matters more than signal existence [27]. Full dimensional decomposition (4D) should dominate single-dimension assessment (partial decomposition), which in turn should dominate direct authorization (no decomposition). This ordering is predicted to hold across domains because the relative information value of decomposed versus holistic assessment is determined by the signal structure, not by corpus-specific features. Isolating the independent contributions of constraint anchoring and role specification requires a factorial design

beyond the scope of this study; the present ablation tests decomposition granularity while holding the other two strategies constant. Managerially, H3 predicts that organizations cannot substitute framing or role instructions for genuine dimensional decomposition; the mechanism's governance effect is carried by the structure of the signal, not by its tone.

### 3. Research Method

#### 3.1. Experimental Design: Within-Subject PA Condition Comparison

Within the design science research tradition, we evaluate the GateFix artifact through a controlled behavioral experiment. The experimental design operationalizes the PA framework through a within-subject comparison of two delegation conditions applied to the same (context, query) pairs: Within-subject experimental designs are widely accepted in behavioral IS research because they control for unobserved task-level heterogeneity by applying both conditions to the same (context, query) pair.

- **Unconstrained condition (Baseline):** The LLM agent receives a standard execution prompt on the (context, query) pair without prior diagnostic requirements. This represents the default delegation relationship in current agentic AI architectures [18, 29], where the principal provides task intent but no structured information constraints. The agent decides whether to proceed based on its own implicit assessment.
- **Constrained condition (Governed):** The same (context, query) pair is first processed through four diagnostic prompts (one per 4D-CQ dimension), and the resulting quality assessment profile is converted to an authorization decision by a deterministic rule. This represents a structured delegation relationship in which the principal imposes pre-execution information constraints that the agent must satisfy before execution is authorized.

The within-subject design ensures that performance differences between conditions are attributable to the governance intervention rather than to variation in task difficulty or context characteristics. The governance-relevant claim is not that LLMs cannot diagnose context quality—they demonstrably can—



but that execution behavior does not by default honor that capacity. To guard against the alternative explanation that observed effects reflect raw differences in prompt length or formatting, both conditions use prompts matched on system-role framing, output-format specification, and token count to within a small tolerance, and both are evaluated on identical (context, query) pairs at identical decoding settings. The only manipulated factor is whether the pre-execution diagnostic stage is present and whether the authorization rule is applied.

All experiments measure LLM judgment under controlled conditions without prior task-specific training, tool calls, or environmental feedback. We operationalize execution at the computational level: producing a numerical or analytical output in response to a query where the output enters an organizational information flow. Extension to advisory and autonomous execution levels is identified as a direction for future work.

### 3.2. Task Corpus and Defect Injection

We use FinQA [10], a benchmark of 8,281 financial question-answer pairs derived from S&P 500 earnings reports with annotated numerical reasoning programs that make execution accuracy verifiable, as the primary task corpus. We construct FinQA-CQ by systematically injecting quality defects along each of the four diagnostic dimensions:

- **Relevance defects** (100 instances): Context replaced with off-topic or wrong-entity financial data
- **Coverage defects** (100 instances): Key numerical quantities removed or made contradictory across sources
- **Ordering defects** (100 instances): Prerequisite quantities placed after dependent computations, temporal ordering made ambiguous
- **Robustness defects** (100 instances): Distractor numbers, OCR-style errors, or adversarial redirections injected into context
- **Clean controls** (100 instances): Unmodified contexts with verified ground truth

Total:  $N = 500$  instances. Inter-annotator agreement on initial coding, before reconciliation: Cohen's  $\kappa = 0.91$ , indicating near-perfect agreement on defect classification. All defect injections were verified by two independent annotators trained on the 4D-CQ dimension definitions; the remaining disagreements were resolved by consensus discussion. The 100-instances-per-dimension sample was chosen to balance annotation cost with the statistical power required to detect the hypothesized main effect ( $\geq 10$  percentage-point improvement in authorization accuracy) at  $p < 0.01$  under a within-subject McNemar test. Observed primary-test effect sizes (43–49 pp for the two above-threshold models) substantially exceeded this conservative planning assumption, yielding ample post-hoc power for the H1 verdict; smaller deltas observed on transfer datasets are interpreted with appropriate caution in Section 4.2. Full prompt templates, annotation guidelines, and defect-injection rules are available as an appendix to support replication.

Each dimension corresponds to a familiar organizational defect pattern. Relevance defects mirror cases where retrieval pulls a similarly named entity rather than the correct one—a sibling subsidiary's filings, or a customer with a near-duplicate name. Coverage defects correspond to source documents missing a required quantity, such as an absent year in a 10-K or an unposted journal entry in a partially synced ledger. Ordering defects match temporal-coherence failures common in batch systems where cause and effect arrive out of sequence. Robustness defects reflect adversarially or accidentally embedded distractors, including OCR scarring on scanned receipts and prompt-injection payloads inside ingested email. The 100-instances-per-dimension scale was chosen to be representative of the labeled samples a typical engineering or risk team can produce in a one- to two-week annotation sprint, balancing statistical power against organizational realism.

For cross-domain validation, we apply the same governance mechanism without recalibration to  $N = 500$  instances each from TAT-QA [11] (hybrid tabular-textual financial QA) and ConvFinQA [12] (multi-turn conversational financial QA). These datasets represent increasing construal levels: FinQA (single-turn, structured, low construal level)  $<$  TAT-QA (cross-format inference, moderate construal level)  $<$  ConvFinQA (multi-turn conversational reasoning, higher construal level). Preserving the prompts

unchanged tests governance portability under construal-level variation without confounding recalibration effects.

### 3.3. Model Selection

Four commercially deployed LLMs were selected to represent architectural diversity and deployment relevance:

- **Qwen-Plus** and **Qwen-Max**: Instruction-tuned models from the Qwen family, producing clear binary quality verdicts when prompted for categorical assessment
- **Claude 3.5 Sonnet**: Alignment-trained proprietary model from Anthropic, known for nuanced, hedged outputs that express uncertainty through qualifications rather than binary verdicts
- **GPT-4o**: Multi-modal proprietary model from OpenAI, representing the frontier capability class with similar hedging tendencies

Selection criteria were threefold. First, coverage of distinct architectural families and training regimes. Second, inclusion of models spanning a range of assessment styles, from binary-preferring to hedged probabilistic outputs. Third, commercial availability at the time of evaluation, so that findings apply directly to the platforms organizations currently use to deploy LLM-based execution pipelines. The exclusion of open-weight architectures such as Llama-3 and Mixtral is a deliberate scoping choice. Because the mechanism depends on a diagnostic-capability threshold rather than on any commercial-API-specific feature, the same logic should plausibly extend to open-weight models—including fine-tuned variants—that meet the threshold, but this extension requires direct empirical replication on open-weight architectures; we revisit this as a generalizability question in Section 6.

### 3.4. Governance Mechanism Operationalization (4D-CQ)

The four-dimensional context quality framework operationalizes Wang and Strong's [30] information quality taxonomy for LLM execution contexts. It extends the original taxonomy in three ways the latter did not anticipate. First, the consumer is an LLM agent with fundamentally different processing

characteristics than a human consumer—position bias, adversarial vulnerability, and no pragmatic inference capacity. Second, consumption occurs in a multi-step pipeline in which diagnostic and authorization functions are sequentially dependent. Third, quality failure triggers automated execution rather than a suboptimal human decision, raising the stakes of misclassification. Full prompt templates for all four dimensions, the deterministic parsing rule, and a worked example are provided in Appendix A to support replication.

The four dimensions, their information quality origins, authorization criteria, and criticality classifications are presented in Table II.

**TABLE II: 4D Context Quality Dimensions and Authorization Mapping**

| <b>Dimension</b> | <b>Information Quality Origin</b>                             | <b>Authorization Criterion</b>                                                                                                                                                                    | <b>Criticality</b>                  |
|------------------|---------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------|
| Relevance        | Contextual quality (relevancy) [30]                           | Context is on-topic for the target task. DEFECT if the context is predominantly off-topic, references the wrong entity or period, or contains more than 30% unrelated material.                   | Critical<br>(REFUSE on DEFECT)      |
| Coverage         | Contextual quality (completeness) [30]                        | Required quantities are present and mutually consistent. DEFECT if any key quantity is absent, contradictory across sources by more than 5%, or relies on estimates rather than reported figures. | Critical<br>(REFUSE on DEFECT)      |
| Ordering         | Representational quality (consistency) [30]                   | Evidence is presented in a logically valid sequence. DEFECT if prerequisite quantities appear after they are used in the computation or temporal ordering is ambiguous.                           | Non-critical<br>(CLARIFY on DEFECT) |
| Robustness       | Intrinsic quality (accuracy / freedom from manipulation) [30] | Context is free from adversarial noise. DEFECT if injected distractors, OCR errors affecting key figures, or sentences designed to redirect the computation are present.                          | Non-critical<br>(CLARIFY on DEFECT) |

- **Authorization rule:** PASS if all four dimensions are assessed as OK; REFUSE if any critical dimension (Relevance or Coverage) returns DEFECT; otherwise CLARIFY (at least one non-critical dimension fails). Table II and Figure 2 give the full mapping; the complete diagnostic prompts and parsing rule are in Appendix A.




Figure 2. Authorization logic flowchart.

Figure 2 presents the authorization logic as a flowchart. The four quality dimensions are evaluated, and failures on critical dimensions trigger execution refusal; failures only on non-critical dimensions trigger a clarification request; passing all dimensions authorizes execution.

### 3.5. Metrics

We define four primary metrics:

- **Diagnostic Accuracy (DimAcc):** Proportion of instances where the LLM's four-dimensional quality assessment matches ground truth across all four dimensions simultaneously.

- **Authorization Accuracy (GateAcc):** Proportion of instances where the final authorization decision (PASS/CLARIFY/REFUSE) matches the correct decision given ground-truth defect status.
- **False PASS Rate (FP):** Proportion of defective contexts incorrectly authorized for execution—the primary safety metric, as it leads to execution on defective inputs.
- **Missed PASS Rate (MP):** Proportion of adequate contexts incorrectly blocked from execution—the primary efficiency metric representing throughput cost, not safety risk.

Statistical significance is assessed via McNemar's test ( $p < 0.01$ ) for within-subject comparisons between the unconstrained and constrained conditions; McNemar's test is the appropriate choice because the primary outcomes are paired binary measurements on the same (context, query) pair across conditions. Given the large effect sizes (authorization-accuracy gains of 10.6–48.8 percentage points and false-pass reductions from approximately 50% to under 5% for models above the diagnostic-capability threshold), the substantive conclusions do not depend on marginal p-value boundaries, even without a formal multiple-comparison correction.

## 4. Results

### 4.1. Main Effect: Governance Intervention Effectiveness (H1)

**TABLE III: Pre-Execution Governance Effectiveness (FinQA-CQ, N = 500)**

Managerial reading: governance raises GateAcc by 30+ percentage points for the strongest models and moves both Qwen models into what we label the operational reliability zone ( $\text{GateAcc} \geq 80\%$ ,  $\text{FP} \leq 5\%$ ), with negligible loss in true-positive execution. We do not claim this zone as a universal standard; it is a laboratory-derived candidate threshold for the model–task combinations studied here, and I&M readers should not interpret it as an industry benchmark without live-system validation.

| Model | DimAcc<br>(Diagnostic) | Baseline<br>GateAcc | Post-<br>Governance<br>GateAcc | False PASS<br>(Post) | Missed PASS<br>(Post) |
|-------|------------------------|---------------------|--------------------------------|----------------------|-----------------------|
|       |                        |                     |                                |                      |                       |

|                      |       |       |       |       |       |
|----------------------|-------|-------|-------|-------|-------|
| Qwen-Plus            | 92.2% | 50.8% | 93.8% | 2.6%  | 3.6%  |
| Qwen-Max             | 71.0% | 36.0% | 84.8% | 4.8%  | 5.2%  |
| Claude 3.5<br>Sonnet | 43.2% | 37.2% | 47.8% | 18.8% | 11.4% |
| GPT-4o               | 27.0% | 40.0% | 56.0% | 23.4% | 8.2%  |




Figure 3. Governance effectiveness across four LLMs (FinQA-CQ, N = 500).

Table III and Figure 3 present the main results. Under unconstrained delegation, all four LLMs achieve authorization accuracy only marginally above the level expected if the execution decision were independent of ground-truth context quality, establishing the baseline against which the governed condition is evaluated in H1.

With the 4D-CQ intervention active, authorization accuracy improves for all four models. Qwen-Plus improves from 50.8% to 93.8% (+43.0 percentage points; McNemar  $p < 0.001$ ). Qwen-Max improves from 36.0% to 84.8% (+48.8 pp;  $p < 0.001$ ). Claude 3.5 Sonnet improves from 37.2% to 47.8% (+10.6

pp;  $p < 0.01$ ), and GPT-4o from 40.0% to 56.0% (+16.0 pp;  $p < 0.01$ ). All improvements are statistically significant.

H1 is supported: the 4D-CQ intervention significantly reduces behavioral biases in LLM execution, with authorization accuracy gains ranging from 10.6 to 48.8 percentage points across the four models tested.

One managerial caveat is worth flagging before the formal boundary analysis. The two models below the laboratory-derived 70% diagnostic-accuracy floor—Claude 3.5 Sonnet and GPT-4o—improve under governance but retain false-pass rates far above the 5% operational ceiling. On those models, GateFix should not be used as the sole safeguard in high-stakes workflows. The full boundary analysis follows in Section 4.4.

The results reveal a clear relationship between diagnostic accuracy and mechanism performance. Qwen-Plus (DimAcc = 92.2%) and Qwen-Max (DimAcc = 71.0%) both enter what we term the reliability zone ( $\text{GateAcc} \geq 80\%$ ,  $\text{FP} \leq 5\%$ ), while Claude 3.5 Sonnet (DimAcc = 43.2%) and GPT-4o (DimAcc = 27.0%) show meaningful but insufficient improvement to reach operational reliability. The pattern suggests a laboratory-derived candidate threshold of approximately 70% diagnostic accuracy—below which the governance mechanism produces unreliable quality signals. We analyze this limiting condition in Section 4.4.

A worked example illustrates the governance mechanism on a concrete instance: a Goldman Sachs revenue growth query (full diagnostic outputs and the four-dimensional verdict are provided in Appendix A.7). Without governance, the LLM correctly diagnoses that key revenue data is missing and sources conflict, yet proceeds to compute an answer—returning +6.0% against a ground truth of +14.1% (an 8.1 percentage-point error, as introduced in the opening vignette). With governance, the quality assessment detects a critical Coverage defect and blocks execution before any computation occurs.

## 4.2. Moderation: Task Construal Level (H2)

**TABLE IV: Cross-Domain Governance Transfer (N = 500 per dataset, Qwen-Plus and Qwen-Max above 70% diagnostic threshold)**



For organizational readers, the practical message is directional: tasks that resemble FinQA—structured, single-turn, with verifiable numerical answers—are the setting in which GateFix delivers its largest accuracy gains; tasks that resemble ConvFinQA—multi-turn, conversational, cross-format—are the setting in which organizations should expect primarily safety improvements rather than large accuracy gains.

| <b>Model</b>      | <b>FinQA <math>\Delta</math> GateAcc</b> | <b>TAT-QA <math>\Delta</math> GateAcc</b> | <b>ConvFinQA <math>\Delta</math> GateAcc</b> |
|-------------------|------------------------------------------|-------------------------------------------|----------------------------------------------|
| Qwen-Plus         | +43.0 pp                                 | +1.8 pp                                   | +5.6 pp                                      |
| Qwen-Max          | +48.8 pp                                 | +3.4 pp                                   | +3.6 pp                                      |
| Claude 3.5 Sonnet | +10.6 pp                                 | -5.6 pp                                   | -2.2 pp                                      |
| GPT-4o            | +16.0 pp                                 | -7.8 pp                                   | -1.4 pp                                      |

Table IV presents the cross-domain transfer results. For Qwen-Plus, the accuracy improvement drops from +43.0 pp on FinQA to +1.8 pp on TAT-QA and +5.6 pp on ConvFinQA. Qwen-Max shows a similar pattern: +48.8 pp on FinQA versus +3.4 pp and +3.6 pp on the transfer datasets. The small positive deltas on transfer datasets should be interpreted with caution: at these effect sizes, statistical power is limited, and the gains may not reach significance in a McNemar test at  $p < 0.01$ . For Claude 3.5 Sonnet and GPT-4o, governance reduces authorization accuracy on the transfer datasets (negative delta values), reflecting the compounding of below-threshold diagnostic accuracy with increased task abstraction.

This attenuation across corpora is consistent with CLT moderation. FinQA tasks are highly structured single-turn computations where the 4D diagnostic dimensions map directly onto concrete, verifiable features—low construal level. TAT-QA introduces hybrid tabular-textual formats requiring cross-format inference, and ConvFinQA requires multi-turn conversational reasoning—both representing higher abstraction levels where the diagnostic decomposition is less precise because the quality dimensions themselves are harder to evaluate in binary terms.

This attenuation is not a design failure but evidence for CLT moderation: when the underlying task is already at a high construal level, decomposition creates less distance between conditions.

**H2 is supported:** on FinQA (the primary test) for models above the laboratory-derived 70% diagnostic-accuracy threshold. On TAT-QA and ConvFinQA, models above the threshold continue to improve, while models below it exhibit attenuation or degradation; this pattern indicates that diagnostic capability, not task abstraction, is the binding constraint. Within the above-threshold cohort, the 4D-CQ intervention shows substantially larger effects for well-structured tasks than for abstract, multi-format tasks, consistent with CLT moderation.

A critical asymmetry emerges in the cross-domain results. While accuracy gains attenuate sharply, the safety guarantee—suppression of false-pass errors—transfers robustly for models above the diagnostic threshold. Qwen-Plus and Qwen-Max maintain FP rates at or under 5% on both TAT-QA and ConvFinQA, despite the substantial accuracy attenuation. This safety–accuracy separability is a core empirical finding: execution safety and execution accuracy are governed by different properties of the diagnostic signal, and they have different portability characteristics, a point developed in Section 5.3.

### 4.3. Decomposition Granularity: Ablation Evidence (H3)

**TABLE V: Ablation Study—Decomposition Granularity (FinQA-CQ, N = 500, Qwen-Plus)**

| Configuration                     | GateAcc | False PASS |
|-----------------------------------|---------|------------|
| Oracle (perfect diagnosis)        | 100%    | 0%         |
| Direct-Auth (no diagnostic stage) | 53.4%   | —          |
| Single-Dim (one composite score)  | 71.2%   | —          |
| Full 4D GateFix                   | 93.8%   | 2.6%       |

Table V presents the ablation analysis testing H3, comparing three governance configurations against an oracle upper bound:

- **Oracle** (perfect diagnosis): Establishes the theoretical ceiling at 100% GateAcc. The governance layer's performance ceiling is determined entirely by the quality of the underlying diagnostic signal.
- **Direct Authorization** (no task decomposition): The model makes a single authorization decision without prior quality decomposition—constraint anchoring and role specification without CLT-motivated dimensional decomposition. Result: 53.4% GateAcc, barely above the unconstrained baseline (50.8%), confirming that specifying evaluation criteria and an evaluator role without decomposing the assessment task provides negligible governance value.
- **Single Dimension** (partial decomposition): A single holistic quality judgment rather than four specific assessments. Result: 71.2% GateAcc—a 17.8 pp improvement over direct authorization, confirming that even partial decomposition provides meaningful value.
- **Full 4D** (complete decomposition): The full four-dimensional mechanism. Result: 93.8% GateAcc—a 22.6 pp improvement over single-dimension assessment and a 40.4 pp improvement over direct authorization.

The monotonic improvement from direct-auth (53.4%) to single-dim (71.2%) to full 4D (93.8%) confirms that task decomposition—the CLT-motivated strategy of lowering abstraction level through dimensional separation—is the primary driver of governance gains. The near-zero incremental value of direct authorization over baseline is theoretically informative: it indicates that constraint anchoring and role specification alone—without decomposition—provide negligible governance value. This is consistent with information design theory's prediction that signal precision matters more than signal existence [27].

**H3 is supported:** finer-grained decomposition yields monotonically greater bias correction (53.4% → 71.2% → 93.8%), and this ordering is preserved as a structural property of the signal design. Isolating the independent contributions of constraint anchoring and role specification remains a direction for future factorial experimentation.

The multi-dimensional structure of GateFix is not merely a design choice—it is necessitated by the multi-dimensional nature of the underlying quality problem. A single holistic quality judgment conflates distinct failure modes (missing data, wrong entity, ordering ambiguity, adversarial noise) into an

undifferentiated signal. Dimensional decomposition allows the authorization function to make differentiated responses—refusing on critical failures while requesting clarification on non-critical ones—resulting in both higher accuracy and lower false-pass rates.

#### **4.4. Boundary Condition: Model Capability Threshold**

The divergence between Qwen-series and Claude/GPT-4o models reveals a prerequisite for the 4D-CQ governance mechanism. Claude 3.5 Sonnet (DimAcc = 43.2%) and GPT-4o (DimAcc = 27.0%) produce qualified, hedged quality assessments rather than the clear binary verdicts the authorization function requires. When parsed into binary decisions, these hedged assessments generate unreliable signals—elevated false-pass rates of 18.8% and 23.4% respectively.

The low diagnostic accuracy of models below this threshold is not a failure of the governance framework; it is a model-capability prerequisite that the framework correctly surfaces. The mechanism accurately reflects the quality of the underlying diagnostic signal rather than masking diagnostic limitations.

The data suggest a laboratory-derived candidate threshold of approximately 70% diagnostic accuracy, below which no deterministic authorization rule can suppress false-pass rates to operationally tolerable levels ( $\leq 5\%$ ). This threshold is consistent with calibration theory's finding that classifier reliability degrades non-linearly below minimum accuracy thresholds [15]. We emphasize that the 70% figure is an observed operational approximation for the model–task combinations studied here, not a theoretical constant. Organizations adopting pre-execution governance in other domains or on other model families should re-estimate the operative threshold on their own data, and should treat 70% as a conservative starting floor rather than a transferable cutoff.




Figure 4. Sensitivity of governance effectiveness to diagnostic-accuracy threshold (Qwen-Plus).




Figure 5. Safety-efficiency operating map across four LLMs (post-governance).

Figure 4 confirms governance stability for models above the threshold: within the recommended threshold range, all metrics vary by fewer than 3 percentage points, indicating that the mechanism does not require precise threshold tuning to maintain reliability. Figure 5 maps the safety-efficiency trade-off across all four models. Qwen-Plus and Qwen-Max move from the ungoverned zone into the reliability region ( $FP \leq 5\%$ ,  $MP \leq 15\%$ ) post-governance. Claude 3.5 Sonnet and GPT-4o improve but remain outside the reliability zone, consistent with the diagnostic capability threshold.

Design-principles bridge. Three principles emerge from the boundary analysis and are developed in detail in §5.1: safety-accuracy separability, minimum diagnostic capability (the laboratory-derived 70% floor), and task-structure contingency. Cross-mapping these to the (concreteness  $\times$  diagnostic-capability) plane yields four practical quadrants. (i) High concreteness  $\times$   $\geq 70\%$  capability is the principal target zone, with

large accuracy gains and reliability-zone safety. (ii) Low concreteness  $\times \geq 70\%$  capability is the safety-only zone—GateFix transfers safety but yields little accuracy lift. (iii) High concreteness  $\times$  sub-70% capability is the model-upgrade zone—governance alone is insufficient and the appropriate response is to substitute or fine-tune the model rather than rely on the mechanism. (iv) Low concreteness  $\times$  sub-70% capability lies outside the mechanism's intended scope and should be served by hybrid human–machine governance.

## 5. Discussion

### 5.1. Key Findings and Design Principles

Our findings converge on three design principles for organizations building AI-enabled execution systems.

#### **Design Principle 1 — Safety–accuracy separability.**

The cross-domain results reveal an asymmetry with significant implications for governance design: the safety guarantee (false-pass suppression) transfers robustly across corpus structures, while accuracy gains are corpus-specific. For Qwen-Plus and Qwen-Max, false-pass rates remain at or under 5% on TAT-QA and ConvFinQA despite accuracy improvements dropping from +43–49 pp to +1.8–5.6 pp. This separability arises because safety depends on diagnostic sensitivity (detecting defects when present), while accuracy depends on diagnostic specificity (correctly clearing adequate contexts). Sensitivity is relatively format-invariant; specificity is format-sensitive. The managerial consequence is a staged deployment strategy: organizations can activate the safety guarantee immediately in new domains, then invest in one-time accuracy calibration for each domain as priority dictates. The safety floor is structural and portable; the accuracy ceiling is configurable and domain-specific. In if–then form: if an organization requires auditable safety properties that transfer across task domains, then execution-layer pre-authorization is a reliable lever, whereas accuracy gains should be evaluated domain by domain.

#### **Design Principle 2 — Minimum diagnostic capability.**

Governance effectiveness is bounded by the underlying LLM's ability to produce clear binary quality verdicts. Our data identify a laboratory-derived candidate threshold of approximately 70% diagnostic accuracy, below which no deterministic authorization rule can suppress false-pass rates to operationally tolerable levels. This is consistent with calibration theory's finding that classifier reliability degrades nonlinearly below minimum accuracy thresholds [15]. For organizations, this implies that model selection for governance compatibility is a first-order design decision: candidate models must be evaluated not only on task performance but on diagnostic clarity—the ability to produce unambiguous binary quality assessments when prompted. In if–then form: if a candidate LLM's diagnostic accuracy on the target context dimensions is below the laboratory-derived approximately 70% threshold, then deterministic authorization gating cannot be relied upon and the organization should either upgrade the model or retain human authorization for that pipeline. Organizations should re-estimate this threshold on their own task distribution before adoption: the 70% value is derived from four commercial models on three financial-QA corpora and is not a transferable constant.

### **Design Principle 3 — Task-structure contingency.**

Governance through dimensional decomposition is most effective when the decomposed subtasks map onto concrete, verifiable features. Structurally concrete domains (financial computation, compliance verification, engineering calculation) are predicted to benefit most from the current 4D decomposition approach, while more abstract domains (strategic analysis, legal interpretation, creative synthesis) may require governance strategies that operate at higher construal levels through schema-level rather than feature-level verification. This is a design principle for matching governance strategy to task type, not a limitation to be overcome. In if–then form: if the task domain is structurally concrete with verifiable outputs, then pre-execution dimensional decomposition is expected to yield substantial accuracy gains in addition to safety improvements; if the domain is abstract or judgment-dependent, the mechanism should be deployed primarily for its safety guarantee rather than for accuracy uplift.



## 5.2. Managerial Implications

We draw four sets of managerial implications from the study, oriented toward organizational leaders designing, deploying, and auditing LLM-based execution systems.

### 5.2.1. Implementation guidance.

When to adopt pre-execution governance. Pre-execution governance delivers the clearest value where (i) the target task is structurally concrete with verifiable outputs, (ii) the underlying LLM achieves at least the laboratory-derived candidate threshold of approximately 70% diagnostic accuracy on the target task type, and (iii) the cost of acting on defective context is high. Typical organizational scenarios meeting these conditions include automated financial analysis and reporting, compliance monitoring, audit-support workflows, and pre-trade or pre-execution risk checks. Organizations operating outside these conditions—especially those deploying models below the diagnostic threshold—should either prioritize model selection or supplement pre-execution governance with alternative controls rather than rely on it as a sole safeguard.

How to integrate GateFix into existing agentic pipelines. GateFix is designed to be deployed as a pre-execution middleware component in existing agentic architectures (for example, ahead of the reasoning step in ReAct-style loops, between retrieval and generation in RAG pipelines, or between agent planning and tool invocation in multi-agent systems). Because the mechanism requires no retraining and no change to the underlying model, it can be inserted incrementally: organizations can wrap existing execution prompts with a 4D-CQ gate, version the gate prompts independently, and maintain gate policies through normal IT governance change-control processes. The computational overhead is four additional LLM API calls per execution request (one per CQ dimension, which can be parallelized and cached) plus a deterministic rule evaluation; this is comparable in order of magnitude to pre-trade risk checks in financial market infrastructure, and should be planned for as part of the pipeline's latency budget.

Throughput versus safety trade-offs. The three-state authorization design (PASS / CLARIFY / REFUSE) gives organizations a policy lever for balancing throughput and safety. Criticality classifications (which

dimensions trigger REFUSE versus CLARIFY) can be configured domain by domain according to risk tolerance; CLARIFY routes borderline cases to human-in-the-loop review without abandoning the task. Our results give managers concrete reference numbers: for capable models, GateFix pushes false-pass rates under 5% at the cost of a missed-pass (throughput) rate in the low-teens—a trade-off most high-stakes workflows would accept, but an explicit one that should be recalibrated for each deployment.

### 5.2.2. Monitoring and fairness.

**Audit readiness and compliance posture.** Because GateFix implements authorization as a deterministic rule over a structured diagnostic signal, it naturally produces an audit trail: every execution decision is accompanied by a four-dimensional quality assessment and a recorded authorization outcome. This aligns with the broader direction of AI governance frameworks that require organizations to document how AI systems are monitored and controlled. Pre-execution governance also instantiates the long-standing IT governance principle of separating policy from mechanism [3, 4, 5] at the inference layer: policy is owned by the organization (which dimensions matter, which trigger REFUSE, what clarifying workflow applies), while the mechanism enforces it deterministically at the agent's execution boundary.

**Human-layer vs. machine-layer governance.** Cognitive forcing functions, the best-tested human-layer intervention, reduce over-reliance by about 30 pp [6]; our machine-layer mechanism produces 43–49 pp gains for capable models. The implication is structural rather than behavioral: human oversight should serve as backstop, structural execution constraints as primary first-line enforcement.

### **Implementation roadmap and monitoring.**

Organizations considering execution-layer governance for LLM agents can follow a four-stage roadmap, independent of the specific LLM vendor. These stages are ordered from pilot feasibility to production discipline, and the last is meant to be continuous.

1. Identify candidate pipelines. (IT-business owners and risk officers nominate; AI governance committee approves the shortlist.) Prioritize LLM-agent workflows that are high-stakes, structurally concrete, and have verifiable outputs—automated financial analysis, compliance checks, audit-support tasks, and pre-

execution risk screens are typical early targets. These are the pipelines in which structural routing failure is most costly and where GateFix is most likely to generate operationally meaningful lift.

2. Configure 4D-CQ for the domain. (Domain SMEs and the ML/AI team co-configure; one to two weeks of labeled data is typically sufficient.) Translate Relevance, Coverage, Ordering, and Robustness into domain-specific decision criteria, and calibrate the REFUSE / CLARIFY thresholds against a small labeled sample of known-good and known-defective contexts. This is the main locus of domain expertise in the deployment; the rest of the harness is largely domain-agnostic.

3. Deploy as an authorization micro-service. (Platform engineering owns the service; the AI governance committee owns gate-policy change control.) Insert GateFix as a pre-execution middleware component between context assembly or retrieval and tool invocation, so that no downstream action can occur without an explicit authorization verdict. Treating authorization as a separate service makes it independently versioned, independently auditable, and substitutable as models improve.

4. Monitor and audit continuously. (Risk and audit functions own the dashboard; the gate log feeds existing IT-risk reporting and internal-audit sampling.) Track false-pass and missed-pass rates, the distribution of PASS / CLARIFY / REFUSE verdicts, and, critically, how these distributions vary across business lines, customer segments, and geographies, so that systematic shifts surface early. The deterministic authorization log is the artifact that connects execution-layer governance to the organization's policy-level AI governance obligations.

Fairness caveat. A higher REFUSE / CLARIFY rate in one segment than another is not by itself evidence that GateFix is fair or unfair; it can reflect genuinely lower context quality upstream. The appropriate response is to investigate and improve data and process quality in that segment, not to raise thresholds to equalize pass rates, because the latter would simply hide a structural disparity behind a uniform throughput number.

### **Applicability beyond financial reasoning.**

Although our experiments focus on financial reasoning, the 4D-CQ structure is designed to map onto any domain where context fitness for execution can be characterized along relevance, coverage, ordering, and

robustness. Two non-financial scenarios illustrate the intended breadth. In clinical decision support, an LLM agent proposing a treatment recommendation operates on a context that should be relevant to the patient in question (not a similarly named record), cover the medications and allergies required for the recommendation, present timeline events in the correct order (diagnosis before intervention, not the reverse), and be robust to transcription errors in the ingested notes. In supply chain anomaly detection, an LLM agent flagging a potential stock-out operates on a context that should be relevant to the affected SKU and lane, cover both inventory and in-transit quantities, preserve event ordering across warehouse and shipping systems, and be robust to partial data ingestion during ERP synchronization windows. In both domains, a GateFix-style authorization layer would block execution on defective context and route borderline cases for human review rather than silently acting. Characterizing the domain-specific thresholds and the appropriate criticality classifications remains a priority for field-study follow-up.

### **Hybrid governance with human oversight.**

The CLARIFY branch of the authorization rule is, by design, the interface between execution-layer structural governance and human oversight. How best to integrate the two remains an open question: the cognitive load on reviewers, the order in which diagnostic dimensions are presented to them, and whether human reviewers should see the agent's own diagnostic signals before forming an independent judgment (a cognitive forcing function [6]) all plausibly shape the quality of the composite decision. We treat the design of hybrid governance architectures—structural constraints at the execution layer combined with cognitive-forcing interventions at the human layer—as a priority research agenda for future work. Two concrete research questions mark this agenda. (RQ-A) Do reviewers who see the agent's own 4D-CQ diagnostic signal before forming a judgment reach more accurate and better-calibrated authorization decisions than reviewers who see only the raw context, and under what cognitive-load conditions does that advantage reverse? (RQ-B) How should the order and granularity in which the four diagnostic dimensions are presented to reviewers be structured to surface systematic upstream data-quality problems rather than merely resolve individual cases?

### 5.3. Theoretical Contributions and Relation to Prior Work

PA theory extension from contract layer to execution layer. Prior work on LLM delegation has made important progress at the contract layer, showing that LLM delegation requires purpose-built mechanism design [7] and that agent information structure shapes behavioral outcomes [8]. Our study adds the execution layer: given a delegation contract, how should the principal structure the signal environment to ensure that the agent's execution behavior is constrained by verifiable quality assessments? The two layers are complementary and jointly necessary. Contract-layer governance specifies what the agent should do and how compliance is verified (*ex ante* structure). Execution-layer governance ensures that the agent's behavior at the moment of action is constrained by its own diagnostic capacity (*ex tempore* structure). A well-designed contract without execution constraints produces an agent that knows its obligations but may not honor them in practice—the behavioral gap our data document.

Evaluation along design-science validity dimensions. Following Hevner et al.'s guidance on design-science evaluation, the three validity dimensions are addressed as follows. Utility is evidenced by the main-effect results: GateFix raises authorization accuracy by 10.6–48.8 percentage points across four commercial models and moves the strongest models into the operational reliability zone ( $FP \leq 5\%$ , Table III). Quality—that the observed effect is driven by the claimed mechanism rather than incidental framing—is evidenced by the ablation analysis in 4.3, which attributes the governance gain to 4D dimensional decomposition rather than to constraint anchoring or role specification. Generalizability is probed through cross-domain evaluation across three financial corpora with different construal-level profiles, yielding the safety–accuracy separability finding and the approximately 70% diagnostic-capability threshold as explicit boundary conditions on transfer.

Structural routing failure as a new governance failure mode. We introduce structural routing failure as a distinctive failure mode—neither classical moral hazard nor adverse selection—in which the agent's diagnostic capacity is accurate and accessible, yet the execution pathway does not route through that assessment by default. The governance remedy is not incentive alignment or monitoring but information

design [27]: restructuring the signal environment so that the agent's diagnostic output is structurally bound to its execution decision. This failure mode is not limited to LLMs. It has structural parallels in human automation bias, where accurate perception of AI unreliability nonetheless fails to translate into appropriate caution [31]. The contrast is instructive: human-layer interventions achieve limited governance gains because they operate indirectly through human behavioral response [6], whereas machine-layer structural constraints can be enforced deterministically. This points toward an architectural implication for IS governance design—human oversight as backstop, structural mechanisms as primary enforcement—that reverses the prevailing design pattern in many current deployments. By conceptualizing an execution harness as an IS governance artifact, this study connects emerging engineering practices around AI harnesses to long-standing IT governance concerns central to Information & Management—control, accountability, and human–automation interaction—thereby placing LLM-agent pipeline design within a recognizable IS research tradition.

Comparison with alternative approaches. GateFix is complementary to, rather than a substitute for, prior techniques. Chain-of-thought prompting [29] improves reasoning quality but operates within the execution pathway and does not structurally separate quality assessment from execution authorization. Retrieval-augmented generation improves context quality but assumes that better context automatically yields better execution decisions—an assumption our data falsify. Fine-tuning and alignment training [32] improve general behavioral tendencies at training time but cannot adapt to deployment-specific governance requirements. GateFix operates at inference time and can be configured and audited without model modification. The ablation evidence in Section 4.3 further shows that the governance effect is not reducible to prompt engineering: direct authorization with constraint anchoring and role specification alone, without 4D decomposition, adds negligible value over the unconstrained baseline (53.4% vs. 50.8% GateAcc). GateFix is therefore a theory-driven information-design artifact in which dimensional decomposition and deterministic aggregation—not framing or role—carry the governance effect, a divergence from raw model behavior that is consistent with dual-process theory [33].

## 6. Limitations and Future Directions

We organize limitations into two tiers. First-order limitations bound the substantive contribution—domain scope (financial computation only), execution type (computational rather than advisory or autonomous), and the laboratory-derived 70% threshold awaiting live-system validation. Second-order limitations bound the empirical generalization—model coverage (no open-weight models), absence of organizational field validation, and ethical/fairness considerations. The remainder of this section addresses each in turn.

First, domain scope and causal identification of the CLT moderator. All experiments are conducted within the financial domain, where FinQA, TAT-QA, and ConvFinQA span a partially ordered range of construal levels. A natural next step is to manipulate construal level within a single dataset—for example, by varying the abstraction level of financial queries on identical underlying data—to deliver a stronger causal test of the moderating mechanism, alongside cross-domain replication in healthcare, supply chain, and compliance settings.

Second, execution type. The experiments operationalize execution at the computational level—producing numerical outputs from structured data. Extension to advisory execution (generating recommendations that enter decision pipelines) and autonomous execution (triggering irreversible actions without human intermediation) constitutes a primary direction for future work. These execution types introduce additional governance complexities—consequence irreversibility, liability attribution, multi-stakeholder accountability—that may require extensions to the 4D framework.

Third, threshold external validity. The approximately 70% diagnostic capability threshold is empirically derived from laboratory conditions and requires live-system validation. In production environments, factors such as context length variation, API latency, model version updates, and adversarial prompt injection may shift the threshold. Continuous monitoring of diagnostic accuracy is a deployment prerequisite.

Fourth, model coverage and mechanism generality. The evaluation uses four commercial LLMs, chosen to mirror the platforms on which organizations currently deploy production LLM pipelines; this decision prioritizes direct managerial relevance but limits generalizability to the open-weight ecosystem. Future work should replicate the mechanism on open-weight models and systematically compare it against prompt-engineering baselines (structured output formatting, system-prompt role assignment, few-shot classification) to isolate the contribution of its theoretical design principles beyond what generic prompt optimization achieves. Given the rapid adoption of open-weight models in enterprise environments, validating GateFix on such architectures is a necessary next step before the mechanism can be codified into organizational standards.

Beyond open-weight coverage, the deterministic parsing rule is sensitive to model output style. Two of the four evaluated models—Claude 3.5 Sonnet and GPT-4o—operated below the 70% diagnostic capability threshold (DimAcc 43.2% and 27.0%, respectively), and we accordingly do not claim safety improvements for these configurations. Two factors plausibly contribute, and we resist attributing the result to a single cause. First, both models exhibit a hedging style, sometimes producing qualified phrasings rather than the binary OK/DEFECT verdict the parser is designed to detect; the parser falls back to DEFECT when no decisive verdict is found (Appendix A.5), depressing measured DimAcc on instances where the model's substantive judgement is in fact correct. Second, even net of parsing variance, diagnostic capability on this task class appears heterogeneous across model families. We interpret this pattern as external evidence supporting Design Principle 2 (minimum diagnostic capability). A concrete future-work direction is the development of parsing rules robust to hedged outputs, which would partially decouple measured diagnostic accuracy from output style and clarify the relative contribution of each factor.

Fifth, organizational field validation. Interaction with existing IT governance structures [3, 4, 5] requires organizational field studies in production deployment settings. The current controlled experimental design establishes mechanism efficacy but does not address organizational adoption dynamics, user trust calibration, workflow integration costs, or interaction between pre-execution constraints and existing



compliance infrastructure. Future field studies in I&M-relevant organizational settings such as financial reporting, compliance monitoring, and customer-facing advisory services can test whether the CLT-moderated effects documented here persist under organizational constraints on latency, cost, and accountability.

Sixth, ethical and fairness considerations merit attention. The mechanism may systematically disadvantage execution on contexts from emerging markets or underrepresented populations where data quality is structurally lower, raising fairness concerns. The REFUSE decision creates liability questions about accountability for blocked execution. And the reliance on binary quality verdicts may suppress contextual nuances that human judgment would preserve. These implications warrant investigation as governance mechanisms move from laboratory conditions to deployment.

Three directions warrant particular emphasis for the I&M research community. (a) Field experiments inserting GateFix into live financial-reporting or compliance pipelines, measuring downstream error rates, throughput, and audit burden simultaneously. (b) Hybrid human-machine governance designs that operationalize the CLARIFY branch as a structured human-review interface, including how the agent's diagnostic signals should be presented to reviewers without inducing automation bias. (c) Cross-domain replication studies in healthcare, supply chain, and customer-facing advisory services, which together test whether the laboratory-derived 70% threshold and the safety-accuracy separability finding hold under live organizational constraints.

## **7. Conclusion**

Our findings show that for LLM execution agents, knowing that the context is defective does not automatically translate into safe behavior; only when diagnostic knowing is structurally tied to doing through an execution harness do organizations obtain reliable safety gains. This study identifies, measures, and addresses a systematic governance failure in LLM-based execution systems: LLMs that accurately diagnose context quality defects nonetheless proceed to execute on defective contexts when

delegation is unconstrained. We frame this failure through PA theory as an asymmetry problem arising from prompt ambiguity, and propose GateFix—a four-dimensional context quality framework with deterministic authorization—as a governance intervention at the execution layer.

The empirical results support all three hypotheses. The 4D-CQ intervention raises authorization accuracy by 10.6 to 48.8 percentage points across four commercial models for models above the approximately 70% diagnostic-accuracy threshold (Table III). Gains are largest on structurally concrete tasks (FinQA: up to +48.8 pp) and attenuate on abstract, multi-format tasks (TAT-QA, ConvFinQA)—a pattern consistent with CLT (Table IV)—while finer-grained decomposition yields monotonically greater bias correction (Table V: 53.4%  $\rightarrow$  71.2%  $\rightarrow$  93.8%).

The study makes three contributions to IS research. Theoretically, it extends PA theory to the execution layer of LLM-based systems—complementing contract-layer work [7, 8]—and introduces CLT [26] as a moderator of intervention effects in human-AI delegation. Methodologically, it contributes the within-subject PA condition comparison design for LLM behavioral experiments and a multi-model, multi-corpus robustness evaluation framework that distinguishes safety portability from accuracy portability. Within the design science research tradition [23, 28], the 4D-CQ mechanism constitutes an instantiated artifact with empirically grounded design principles. Practically, it provides a deployable governance mechanism with empirically derived deployment prerequisites: a 70% diagnostic accuracy threshold and construal-level-dependent effectiveness that together determine the mechanism's domain of reliable applicability.

Two findings carry particular significance for the design of governed AI systems. The safety–accuracy separability principle—the discovery that execution safety and accuracy are separable governance properties with different transfer characteristics—suggests that organizations can deploy safety-first governance immediately while iteratively calibrating for accuracy in each target domain. The CLT moderation finding provides a theoretically grounded decision framework for determining where pre-execution governance delivers the most value: in domains characterized by concrete, verifiable, structured execution tasks.

A complementary contribution grounds an emergent engineering concept in IS theory. The execution harness—a governance environment that constrains and directs LLM execution to ensure reliable outputs—has been independently developed across frontier AI deployments as an architectural response to the structural failure this study documents [22]. By formalizing the harness as an IS governance artifact and grounding it in IT governance theory, information quality theory, and information design theory, this work supports progress toward auditable harness design, cross-institutional deployment standards, and a theoretical basis for regulatory oversight—objectives that complement, rather than substitute for, the engineering practice from which the harness concept emerged.

As organizations scale their deployment of LLM agents, closing the loop between what an agent knows about its context and what it is allowed to do is not an optional refinement—it is a structural prerequisite for trustworthy AI-enabled execution. The question for IS governance research is therefore not how to build AI systems that govern themselves, but how to design the information structures that translate organizational governance intent into enforceable constraints on automated execution. This paper contributes one of the first empirically grounded execution-layer mechanisms for that purpose and establishes the boundary conditions for its effective deployment.

Data and code availability. The FinQA-CQ defect-injection protocol, the 4D-CQ prompt templates, the deterministic authorization rule, and the analysis scripts used in this study will be shared in an appendix and a public repository upon acceptance, subject to the licensing terms of the underlying benchmark corpora (FinQA, TAT-QA, ConvFinQA).

## **Appendix A. Prompt Templates, Parsing Rules, and Worked Example**

This appendix provides the complete prompt templates and deterministic output-parsing rule used for each of the four diagnostic dimensions in GateFix, the authorization mapping from dimension scores to PASS / CLARIFY / REFUSE, and one worked example. All prompts are applied identically across the

four backbone models (Qwen-Plus, Qwen-Max, Claude 3.5 Sonnet, GPT-4o) without per-model tuning, so reproductions on new backbones do not require prompt rewrites.

#### A.1 Relevance dimension (s\_rel)

The Relevance prompt assesses whether the retrieved context is semantically aligned with the task goal and does not contain substantial off-topic material that could mislead numerical reasoning.

SYSTEM: You are a financial context quality evaluator. Answer only "OK" or "DEFECT".

USER: Task goal: {goal}

Retrieved context: {context}

Evaluate whether the retrieved context is relevant to the task goal. A context is DEFECT if it is predominantly off-topic, references the wrong company or time period, or contains more than 30% unrelated material. Otherwise respond OK.

Respond with exactly one word: OK or DEFECT.

#### A.2 Coverage dimension (s\_cov)

The Coverage prompt checks whether the context contains sufficient numerical evidence to execute the requested computation without requiring external assumptions or imputation of missing values.

SYSTEM: You are a financial context quality evaluator. Answer only "OK" or "DEFECT".

USER: Task goal: {goal}

Retrieved context: {context}

Evaluate whether the context provides sufficient numerical evidence to complete the task. A context is DEFECT if any key quantity required for the computation is absent, contradictory across sources by more than 5%, or relies on estimates rather than reported figures. Otherwise respond OK.

Respond with exactly one word: OK or DEFECT.

#### A.3 Ordering dimension (s\_ord)

The Ordering prompt evaluates whether the reasoning steps or evidence presented in the context follow a logically valid sequence that supports the target computation.

SYSTEM: You are a financial context quality evaluator. Answer only "OK" or "DEFECT".

USER: Task goal: {goal}

Retrieved context: {context}

Evaluate whether the evidence in the context is presented in a logically ordered sequence for the task. A context is DEFECT if prerequisite quantities appear after they are used, or if the temporal ordering of financial figures is ambiguous or internally inconsistent. Otherwise respond OK.

Respond with exactly one word: OK or DEFECT.

#### A.4 Robustness dimension (s\_rob)

The Robustness prompt checks whether the context is free from adversarial noise, typographical corruption, or injected distractors that could destabilize numerical reasoning.

SYSTEM: You are a financial context quality evaluator. Answer only "OK" or "DEFECT".

USER: Task goal: {goal}

Retrieved context: {context}

Evaluate whether the context contains noise or adversarial perturbations that could corrupt numerical reasoning. A context is DEFECT if it contains injected distractor numbers, OCR errors affecting key figures, or sentences designed to redirect the computation toward incorrect values. Otherwise respond OK.

Respond with exactly one word: OK or DEFECT.

#### A.5 Output parsing and fallback rule

Raw model outputs are parsed by the following deterministic procedure applied uniformly across all four backbones: (1) strip leading/trailing whitespace and convert to upper case; (2) if the result equals "OK", assign s\_i = OK; (3) if the result equals "DEFECT", assign s\_i = DEFECT; (4) if neither matches (ambiguous output), apply the conservative fallback s\_i = DEFECT. In our experiments, ambiguous outputs accounted for fewer than 1.3% of all invocations across the FinQA-CQ benchmark; no post-processing beyond this rule is applied.

```
def parse_response(raw: str) -> str:
```

```
    token = raw.strip().upper()
```

```

if token == "OK":    return "OK"

if token == "DEFECT": return "DEFECT"

return "DEFECT" # conservative fallback

```

#### A.6 Authorization rule (PASS / CLARIFY / REFUSE)

Given the four parsed dimension scores  $s\_rel, s\_cov, s\_ord, s\_rob \in \{OK, DEFECT\}$ , the deterministic authorization function maps the score vector to one of three outcomes. The two critical dimensions (Relevance, Coverage) trigger REFUSE on any DEFECT; the two non-critical dimensions (Ordering, Robustness) trigger CLARIFY when they are the only sources of failure; PASS requires all four to be OK.

```

def authorize(s_rel, s_cov, s_ord, s_rob) -> str:

    if s_rel == "DEFECT" or s_cov == "DEFECT":

        return "REFUSE"

    if s_ord == "DEFECT" or s_rob == "DEFECT":

        return "CLARIFY"

    return "PASS"

```

#### A.7 Worked example

We illustrate the four-dimension evaluation and authorization rule on the Goldman Sachs revenue-growth instance discussed in Section 1 (Figure 1).

Task goal. Compute Goldman Sachs's year-over-year revenue growth for the most recent reporting period.

Retrieved context (abridged). Two excerpts are concatenated: (i) a 10-K segment containing Goldman Sachs's 2019 net revenues but missing the 2018 line item; (ii) an investor-relations summary citing a 2018 figure that conflicts with a third-party data provider's number by approximately 8%. No off-topic content is present, ordering is internally consistent, and no adversarial perturbations are detected.

Diagnostic outputs. Relevance: OK (context is on the correct entity and period). Coverage: DEFECT (2018 revenue is absent in (i) and inconsistent in (ii) by more than 5%). Ordering: OK. Robustness: OK.

Authorization. `authorize(OK, DEFECT, OK, OK) = REFUSE`, because Coverage is critical. The agent refuses to compute the growth rate and surfaces a structured explanation citing the missing and inconsistent 2018 figure. Without GateFix, the same backbone proceeded to compute a +6.0% growth estimate against an audited ground truth of +14.1%, an 8.1 percentage-point error — the failure mode the paper documents and the mechanism is designed to suppress.

## **Declaration of Generative AI and AI-assisted Technologies in the Writing Process**

During the preparation of this manuscript, the authors used generative AI tools (large language models) for language editing and formatting assistance. After using these tools, the authors reviewed and edited the content as needed and take full responsibility for the content of the publication. Generative AI was not used as a substitute for human critical thinking, research design, data analysis, or interpretation of results.

## **References**

- [1] PromptArmor. "Data Exfiltration from Slack AI via Indirect Prompt Injection." Security analysis report, PromptArmor (blog/Substack), August 2024.
- [2] Microsoft Security Response Center. "CVE-2025-32711: Microsoft 365 Copilot Information Disclosure Vulnerability (EchoLeak)." MSRC Advisory, June 2025.
- [3] Sambamurthy, V., and Zmud, R.W. "Arrangements for information technology governance: A theory of multiple contingencies." *MIS Quarterly*, 23, 2 (1999), 261–290.
- [4] Weill, P., and Ross, J.W. *IT Governance: How Top Performers Manage IT Decision Rights for Superior Results*. Harvard Business School Press, Boston, MA, 2004.
- [5] ISACA. *COBIT 2019 Framework: Governance and Management Objectives*. ISACA, Rolling Meadows, IL, 2019.

- [6] Buçinca, Z.; Malaya, M.B.; and Gajos, K.Z. "To trust or to think: Cognitive forcing functions can reduce overreliance on AI in AI-assisted decision making." *Proceedings of the ACM on Human-Computer Interaction*, 5, CSCW1 (2021), Article 188.
- [7] Kirshner, S.N.; Pan, Y.; and Wu, J.X. "The AI Agent's Dilemma: LLM Contract Design under Moral Hazard." Working Paper, 2025. Available at SSRN: 5607356.
- [8] Kirshner, S.N.; Pan, Y.; Wu, J.X.; and Gould, A. "Talking Terms: Agent Information in LLM Supply Chain Bargaining." *Decision Sciences*, 2026 (forthcoming). Available at SSRN: 5273609.
- [9] Wu, J.X.; Wu, Y.D.; Chen, K.Y.; and Hua, L. "Building socially intelligent AI systems: Evidence from the trust game using artificial agents with deep learning." *Management Science*, 69, 12 (2023), 7236–7252.
- [10] Chen, Z.; Chen, W.; Smiley, C.; Shah, S.; Borova, I.; Langdon, D.; Moussa, R.; Beane, M.; Huang, T.H.; Routledge, B.; and Wang, W.Y. "FinQA: A dataset of numerical reasoning over financial data." In *Proceedings of EMNLP*, 2021, pp. 3697–3711.
- [11] Zhu, F.; Lei, W.; Huang, Y.; Wang, C.; Zhang, S.; Lv, J.; Feng, F.; and Chua, T.S. "TAT-QA: A question answering benchmark on a hybrid of tabular and textual content in finance." In *Proceedings of ACL*, 2021, pp. 3277–3287.
- [12] Chen, Z.; Li, C.; Smiley, C.; Ma, Z.; Shah, S.; and Wang, W.Y. "ConvFinQA: Exploring the chain of numerical reasoning in conversational finance question answering." In *Proceedings of EMNLP*, 2022, pp. 6279–6292.
- [13] Chen, Y.; Kirshner, S.N.; Ovchinnikov, A.; Andiappan, M.; and Jenkin, T. "A manager and an AI walk into a bar: Does ChatGPT make biased decisions like we do?" *Manufacturing and Service Operations Management*, 27 (2025), 354–368.
- [14] Kirshner, S.N.; Pan, Y.; and Wu, J.X. "Prosocial When Simple and Cold-Hearted When Complex: How Task Difficulty Shapes LLM Behavior." *Decision Analysis*, 2025. Available at SSRN: 5266963.
- [15] Guo, C.; Pleiss, G.; Sun, Y.; and Weinberger, K.Q. "On calibration of modern neural networks." In *Proceedings of the 34th ICML*, 2017, pp. 1321–1330.



- [16] Shapira, N.; Wendler, C.; Yen, A.; Sarti, G.; Pal, K.; Floody, O.; et al. "Agents of chaos: Red-teaming autonomous LLM agents in live environments." arXiv preprint arXiv:2602.20021, 2026.
- [17] Liu, N.F.; Lin, K.; Hewitt, J.; Paranjape, A.; Bevilacqua, M.; Petroni, F.; and Liang, P. "Lost in the middle: How language models use long contexts." *Transactions of the Association for Computational Linguistics*, 12 (2024), 157–173.
- [18] Yao, S.; Zhao, J.; Yu, D.; Du, N.; Shafran, I.; Narasimhan, K.; and Cao, Y. "ReAct: Synergizing reasoning and acting in language models." In *Proceedings of the 11th International Conference on Learning Representations*, 2023.
- [19] Wiener, N. *Cybernetics: Or Control and Communication in the Animal and the Machine*. MIT Press, Cambridge, MA, 1948.
- [20] Gartner. "Gartner predicts over 40% of agentic AI projects will be canceled by end of 2027." Gartner Press Release, June 25, 2025.
- [21] Open Worldwide Application Security Project. *OWASP Top 10 for Large Language Model Applications*, Version 2025. OWASP Foundation, 2025.
- [22] Anthropic Engineering. "Harness design for long-running application development." Anthropic Engineering Blog, 24 March 2026. Available at: <https://www.anthropic.com/engineering/harness-design-long-running-apps> (accessed 29 April 2026).
- [23] Hevner, A.R.; March, S.T.; Park, J.; and Ram, S. "Design science in information systems research." *MIS Quarterly*, 28, 1 (2004), 75–105.
- [24] Eisenhardt, K.M. "Agency theory: An assessment and review." *Academy of Management Review*, 14, 1 (1989), 57–74.
- [25] Ross, S.A. "The economic theory of agency: The principal's problem." *American Economic Review*, 63, 2 (1973), 134–139.
- [26] Trope, Y., and Liberman, N. "Construal-level theory of psychological distance." *Psychological Review*, 117, 2 (2010), 440–463.

- [27] Bergemann, D., and Morris, S. "Information design: A unified perspective." *Journal of Economic Literature*, 57, 1 (2019), 44–95.
- [28] Gregor, S., and Hevner, A.R. "Positioning and presenting design science research for maximum impact." *MIS Quarterly*, 37, 2 (2013), 337–355.
- [29] Wei, J.; Wang, X.; Schuurmans, D.; Bosma, M.; Xia, F.; Chi, E.; Le, Q.V.; and Zhou, D. "Chain-of-thought prompting elicits reasoning in large language models." *Advances in Neural Information Processing Systems*, 35 (2022), 24824–24837.
- [30] Wang, R.Y., and Strong, D.M. "Beyond accuracy: What data quality means to data consumers." *Journal of Management Information Systems*, 12, 4 (1996), 5–33.
- [31] Parasuraman, R., and Riley, V. "Humans and automation: Use, misuse, disuse, abuse." *Human Factors*, 39, 2 (1997), 230–253.
- [32] Anthropic. "Constitutional AI: Harmlessness from AI feedback." Technical Report arXiv:2212.08073, 2022.
- [33] Kahneman, D. *Thinking, Fast and Slow*. Farrar, Straus and Giroux, New York, 2011.